

TTS Campaign Episode Processor - Project Outline

Project Overview

A web application that converts D&D campaign episode scripts (markdown format) into high-quality TTS audio using ElevenLabs API, with intelligent speaker detection, batch processing, and cost controls.

Project Structure

campaign-tts-processor/		
├── frontend/		
	├── index.html	# Main interface
	├── app.js	# Core application logic
	├── styles.css	# Clean, focused styling
	└── components/	
	├── upload.js	# File upload handling
	├── parser.js	# Markdown/script parser
	├── preview.js	# Speaker preview component
	└── processor.js	# TTS processing logic
├── backend/		
	├── server.js	# Express server
	├── api/	
	├── elevenlabs.js	# ElevenLabs API wrapper
	├── auth.js	# API key management
	└── queue.js	# Processing queue
	├── utils/	
	├── parser.js	# Server-side parsing
	├── splitter.js	# Text chunking for API limits
	└── merger.js	# Audio file merging
	└── config/	
	└── voices.json	# Voice configuration
├── outputs/		
# Generated audio files		
├── .env.example		
# Environment variables template		
└── README.md		
# Setup instructions		

Core Features (MVP)

1. File Upload & Parsing

javascript

// Parse episode format

```
class EpisodeParser {
  parseScript(markdown) {
    const scenes = [];
    const speakers = new Set();

    // Split by timestamps (e.g., ## **00:47:00**)
    const sections = markdown.split(/##\s*\*\*\d{2}:\d{2}:\d{2}\s*\*/);

    sections.forEach(section => {
      const dialogues = [];

      // Match speaker patterns: **SPEAKER:** or **SPEAKER (DESCRIPTION):**
      const speakerRegex = /\*\*([A-Z\s]+)(?:\s*\([^\)]+\))?\s*\*\*/g;

      let match;
      while ((match = speakerRegex.exec(section))) {
        const speaker = match[1].trim();
        speakers.add(speaker);

        // Extract dialogue until next speaker or section end
        const startIndex = match.index + match[0].length;
        const nextMatch = speakerRegex.exec(section);
        const endIndex = nextMatch ? nextMatch.index : section.length;

        const dialogue = section.substring(startIndex, endIndex).trim();

        // Check if dialogue needs chunking (rare but important for DM narration)
        if (dialogue.length > 1000) {
          console.log(`Long dialogue detected for ${speaker}: ${dialogue.length} char`);
        }

        dialogues.push({ speaker, text: dialogue });

        // Reset regex position
        speakerRegex.lastIndex = nextMatch ? nextMatch.index : section.length;
      }

      scenes.push(dialogues);
    });

    return { scenes, speakers: Array.from(speakers) };
  }
}
```

2. Speaker Detection & Mapping

javascript

```
// Auto-detect all speakers in episode  
const speakers = {  
  "JOE": { voice: "Daniel", settings: {...} },  
  "TRONALD": { voice: "Adam", settings: {...} },  
  "SAXY": { voice: "Bella", settings: {...} },  
  // etc...  
}
```

3. Processing Interface

- Upload episode file
- Auto-detect speakers
- Map speakers to ElevenLabs voices
- Preview each speaker's voice
- Set processing options
- Monitor progress
- Download results

4. Cost Controls

javascript

// Failsafes to prevent overspending


```
const costControls = {
  maxCharactersPerChunk: 1200,    // Optimal for natural speech
  maxCharactersPerRun: 50000,     // Total episode limit
  estimateCost: (characters) => characters * 0.00003,
  requireConfirmation: true,
  dailyLimit: 1000000, // characters
  warningThreshold: 0.8,

  // Auto-chunk detection
  autoChunkThreshold: 1000, // Chunk anything over 1000 chars

  validateChunk: (text) => {
    if (text.length > 4500) {
      throw new Error('Single chunk exceeds ElevenLabs limit! This should never happen.')
    }
    return true;
  }
}
```


Frontend Design


Main Interface Layout


 Campaign TTS Processor


[Upload Episode File]


Detected Speakers:


JOE (DM) → [Daniel ▼] 


TRONALD → [Adam ▼] 

SAXY → [Bella ▼] 

CELESTIA → [Rachel ▼] 


RAARH → [Antoni ▼] 

MITE → [Josh ▼] 

ALDEN → [Sam ▼] 

Episode Info:

- Episode: 117A
- Runtime: ~84 minutes
- Total Dialogue: 112 sections
- Characters: 87,293
- Est. Cost: \$2.62
- Long Sections: 3 (will auto-chunk)

 3 sections over 1000 chars detected


These will be naturally chunked

[Process Episode]

[Advanced Options]

Progress: 45%

Processing: Scene 3 of 7...

 **ElevenLabs Integration**

API Configuration

javascript

// backend/api/elevenLabs.js

```
const ElevenLabs = require('elevenlabs-api');
```

```
class TTSPProcessor {
```

```
  constructor(apiKey) {
```

```
    this.client = new ElevenLabs({ apiKey });
```

```
    this.voiceCache = new Map();
```

```
    this.chunker = new DialogueChunker();
```

```
  }
```

```
  async processDialogue(text, speaker, settings) {
```

```
    // Smart chunking for Long dialogue (mainly DM narration)
```

```
    const chunks = this.chunker.chunkLongDialogue(text, speaker);
```

```
    const audioFiles = [];
```

```
    for (const chunk of chunks) {
```

```
      console.log(`Processing ${speaker} chunk ${chunk.chunkNumber} (${chunk.text.length}
```

```
      const voice = this.voiceCache.get(speaker) ||
```

```
        await this.loadVoice(speaker);
```

```
      const audio = await this.client.textToSpeech({
```

```
        text: chunk.text,
```

```
        voice_id: voice.id,
```

```
        model_id: "eleven_monolingual_v1",
```

```
        voice_settings: {
```

```
          stability: settings.stability || 0.75,
```

```
          similarity_boost: settings.similarity || 0.75,
```

```
          style: settings.style || 0.5,
```

```
          use_speaker_boost: true
```

```
        }
```

```
      });
```

```
      audioFiles.push({
```

```
        speaker,
```

```
        chunkNumber: chunk.chunkNumber,
```

```
        audio,
```

```
        naturalBreak: chunk.naturalBreak
```

```
      });
```

```
    }
```

```
    return audioFiles;
```

```
  }
```

```
}
```

Smart Text Chunking for Natural Speech

// Intelligent chunking for natural TTS flow

```
class DialogueChunker {
  constructor() {
    // ElevenLabs sounds best with 700-1200 chars
    this.optimalChunkSize = 1000;
    this.maxChunkSize = 1200;
    this.minChunkSize = 500;
  }

  chunkLongDialogue(text, speaker) {
    // Most dialogue is under 1000 chars, but DM sections can be longer
    if (text.length < this.maxChunkSize) {
      return [{ text, speaker, chunkNumber: 1 }];
    }

    const chunks = [];
    const paragraphs = text.split(/\n\n+/);
    let currentChunk = '';
    let chunkNumber = 1;

    for (const para of paragraphs) {
      // If adding this paragraph exceeds optimal size, split here
      if (currentChunk &&
        (currentChunk.length + para.length > this.optimalChunkSize)) {

        // Save current chunk
        chunks.push({
          text: currentChunk.trim(),
          speaker,
          chunkNumber: chunkNumber++,
          naturalBreak: true
        });
        currentChunk = para;
      } else {
        // Add to current chunk
        currentChunk += (currentChunk ? '\n\n' : '') + para;
      }

      // Emergency split if single paragraph is too long
      if (currentChunk.length > this.maxChunkSize) {
        const sentences = currentChunk.match(/^[^!?]+[.!?]+/g) || [];
        let tempChunk = '';

        for (const sentence of sentences) {
          if (tempChunk.length + sentence.length > this.optimalChunkSize) {
            chunks.push({
```

```

        text: tempChunk.trim(),
        speaker,
        chunkNumber: chunkNumber++,
        naturalBreak: false
    });
    tempChunk = sentence;
    } else {
        tempChunk += ' ' + sentence;
    }
}
currentChunk = tempChunk;
}
}

// Don't forget the last chunk
if (currentChunk.trim()) {
    chunks.push({
        text: currentChunk.trim(),
        speaker,
        chunkNumber: chunkNumber,
        naturalBreak: true
    });
}

return chunks;
}

// Example output for DM narration:
// Chunk 1: "You step into the courtyard..." (743 chars) ✓
// Chunk 2: "The statues around you..." (891 chars) ✓
// Chunk 3: "As you examine the seal..." (656 chars) ✓
}

```

Real-World Example

javascript

```
// Your transcript rarely has super long sections, but when it does:  
const dmNarration = `You step into what was once a grand courtyard...  
[700 characters of description]
```

The ground is cracked flagstone, weeds pushing through...
[500 characters of description]

Scattered across the area are the remains...
[800 characters of description]`;

```
const chunks = chunker.chunkLongDialogue(dmNarration, 'JOE');  
// Returns 3 natural chunks instead of awkward mid-sentence breaks
```

Security & Error Handling

API Key Management

javascript

```
// Never expose API keys in frontend  
// backend/.env
```

```
ELEVENLABS_API_KEY=your_key_here  
DAILY_CHAR_LIMIT=1000000  
COST_PER_CHAR=0.00003
```

```
// backend/auth.js
```

```
const validateRequest = (req, res, next) => {  
  const { characters } = req.body;  
  const cost = characters * process.env.COST_PER_CHAR;  
  
  if (cost > 50) { // $50 limit  
    return res.status(403).json({  
      error: 'Cost limit exceeded',  
      estimated: cost  
    });  
  }  
  next();  
};
```

Error Recovery

```
javascript
```

```
// Retry failed chunks
async function processWithRetry(chunk, maxRetries = 3) {
  for (let i = 0; i < maxRetries; i++) {
    try {
      return await processChunk(chunk);
    } catch (error) {
      if (i === maxRetries - 1) throw error;
      await delay(1000 * (i + 1)); // Exponential backoff
    }
  }
}
```

Implementation Steps (for Cursor)

Step 1: Basic Setup

```
bash

# Create project structure
mkdir campaign-tts-processor
cd campaign-tts-processor
npm init -y
npm install express dotenv multer
npm install elevenlabs-node-api # or official SDK
```

Step 2: File Upload & Parsing

Tell Cursor:

"Create a file upload interface that accepts markdown files. Parse the markdown to extract speaker names (format: **SPEAKER_NAME:**) and their dialogue. Show me a preview of detected speakers."

Step 3: Voice Mapping Interface

Tell Cursor:

"Create a speaker-to-voice mapping interface. Each detected speaker should have a dropdown to select an ElevenLabs voice and a preview button to test the voice with sample text."

Step 4: Processing Engine

Tell Cursor:

"Create a processing engine that takes the parsed dialogue and sends it to ElevenLabs API. For any dialogue over 1000 characters, use the DialogueChunker class to split at paragraph breaks. Show

which sections are being chunked in the UI. Save audio files with format:
Episode_117A_Scene_1_JOE_chunk1.mp3"

Step 5: Audio Merging

Tell Cursor:

"Create a system to merge all generated audio files in the correct order based on timestamps. Output a single Episode_117A_Complete.mp3 file."

UI/UX Guidelines

Design Principles

- **Clarity First:** Show what's happening at every step
- **Cost Visible:** Always show estimated cost before processing
- **Progress Tracking:** Real-time updates during processing
- **Error Recovery:** Clear error messages with solutions

Color Scheme

```
css

:root {
  --primary: #5b21b6;      /* Purple - main actions */
  --secondary: #1e293b;    /* Dark slate - headers */
  --success: #10b981;      /* Green - completed */
  --warning: #f59e0b;      /* Amber - warnings */
  --danger: #ef4444;       /* Red - errors */
  --bg: #f8fafc;           /* Light background */
  --text: #334155;         /* Main text */
}
```

Advanced Features (Phase 2)

1. Multi-Voice Scenes

- Detect when multiple characters speak together
- Process overlapping dialogue appropriately

2. Sound Effects Integration

- Add ambient sounds between scenes
- Include dice rolling sounds for combat

3. Voice Memory

- Save speaker-to-voice mappings
- Auto-apply to future episodes

4. Batch Processing

- Queue multiple episodes
 - Process overnight with rate limiting
-

Testing Checklist

- ☐ Upload markdown file
 - ☐ Correctly parse all speakers
 - ☐ Map voices and preview
 - ☐ Calculate accurate cost
 - ☐ Process small test section
 - ☐ Handle API errors gracefully
 - ☐ Merge audio correctly
 - ☐ Download final file
-

Cursor Communication Guide

When implementing with Cursor, use these exact phrases:

1. **"Create the basic file structure as shown in the project outline"**
2. **"Implement the episode parser to extract speakers and dialogue from markdown"**
3. **"Build the speaker detection and voice mapping interface"**
4. **"Create the ElevenLabs API integration with proper error handling"**
5. **"Add cost calculation and confirmation before processing"**
6. **"Implement the audio processing queue with progress tracking"**
7. **"Create the audio file merger to combine all segments"**

Remember to test each component individually before connecting them!

Important Notes

- **NEVER** put API keys in frontend code
- **ALWAYS** show cost estimates before processing
- **CHUNK** text properly to avoid API limits
- **SAVE** progress in case of failures

- **TEST** with small sections first

This focused approach gives you exactly what you need without unnecessary complexity. Start with the MVP and add features as needed!